

Exercise 2

Advanced Methods for Regression and Classification

Martina Adamuscinova

2025-10-21

Load and split the data:

The data contains information about construction costs `df$y` of real estate single-family residential apartments in Tehran, Iran.

```
setwd('/Users/martinaadamuscinova')
load("building.RData")

set.seed(24)


n <- nrow(df)
train <- sample(1:n, size = round(2/3 * n))
test <- (1:n)[-train]
```

Exercise 1

```
m1<- lm(y~, data = df, subset = train)
m1_pred_test <- predict(m1, newdata = df[test,])
```

a)

```
par(mfrow = c(2, 2), cex = 0.7, mar = c(4, 4, 2, 2))
plot(m1)
```



From the ‘Residuals vs Fitted’ and ‘Scale–Location’ plots, it seems that the homoscedasticity assumption is satisfied. Residuals are fairly evenly scattered with no clear pattern. However, slight curvature can be detected in the plots which could indicate mild non-linearity. From the ‘Q–Q Residuals’ plot we see that most residuals follow the normal line closely, though deviations occur at both tails, pointing to the potential presence of outliers. The ‘Residuals vs Leverage’ plot indicates a few observations that could influence the regression coefficients, but overall influence seems moderate.

b)

coef(m1)

```

##          (Intercept)      START.YEAR      START.QUARTER    COMPLETION.YEAR
## 3.060275e+02 -4.583304e+00 -6.491958e-01 7.747832e-02
## COMPLETION.QUARTER      PhysFin1      PhysFin2      PhysFin3
## 2.575343e-02 -2.957669e-02 3.202121e-05 7.620888e-05
##          PhysFin4      PhysFin5      PhysFin6      PhysFin7
## -1.062381e-04 -3.112674e-03 6.660530e-04 NA
##          PhysFin8      Econ1       Econ2       Econ3
## 4.867836e-04 6.459288e-04 3.510751e-01 4.535328e-01
##          Econ4       Econ5       Econ6       Econ7
## -2.739181e-02 2.865240e-05 -5.126585e-04 -1.653414e-01
##          Econ8       Econ9      Econ10      Econ11
## 1.145921e-02 -3.146889e-04 9.555793e-01 -2.710021e-03
##          Econ12      Econ13      Econ14      Econ15
##
```

```

##      4.087347e-03    -1.540769e-04     1.200412e-03    -3.096889e-01
##          Econ16           Econ17           Econ18           Econ19
##      9.643135e-01     5.505590e-05     1.578076e-04    -7.579301e-06
##          Econ1.lag1        Econ2.lag1        Econ3.lag1        Econ4.lag1
##      -5.100883e-04    -9.855374e-01    -1.692274e-01     6.684749e-01
##          Econ5.lag1        Econ6.lag1        Econ7.lag1        Econ8.lag1
##      1.299865e-05     8.645901e-04    -5.323517e-02    -6.956040e-03
##          Econ9.lag1        Econ10.lag1       Econ11.lag1       Econ12.lag1
##      5.252813e-05    -5.369477e-01    -5.840291e-03     9.687316e-04
##          Econ13.lag1        Econ14.lag1       Econ15.lag1       Econ16.lag1
##      1.120091e-04     3.085272e-04     9.695643e-01    -1.211819e+00
##          Econ17.lag1        Econ18.lag1       Econ19.lag1       Econ1.lag2
##      -6.210047e-04     8.009613e-05    -9.024392e-06    -2.390565e-04
##          Econ2.lag2        Econ3.lag2        Econ4.lag2       Econ5.lag2
##      1.049877e+00    -5.013525e-02    -4.123327e-01    -1.770217e-05
##          Econ6.lag2        Econ7.lag2        Econ8.lag2       Econ9.lag2
##      -5.325516e-04    -8.112318e-02    -1.355449e-02    -1.145171e-04
##          Econ10.lag2        Econ11.lag2       Econ12.lag2       Econ13.lag2
##      4.462067e-01    -5.264031e-03     1.068559e-02     1.070018e-04
##          Econ14.lag2        Econ15.lag2       Econ16.lag2       Econ17.lag2
##      -6.555360e-04    -3.294641e-01     7.179874e-01    -1.993108e-04
##          Econ18.lag2        Econ19.lag2       Econ1.lag3       Econ2.lag3
##      1.555121e-04     5.096244e-06     8.471132e-04    -7.834777e-01
##          Econ3.lag3        Econ4.lag3        Econ5.lag3       Econ6.lag3
##      1.134360e-01    -2.288266e-01            NA            NA
##          Econ7.lag3        Econ8.lag3       Econ9.lag3       Econ10.lag3
##          NA            NA            NA            NA
##          Econ11.lag3        Econ12.lag3       Econ13.lag3       Econ14.lag3
##          NA            NA            NA            NA
##          Econ15.lag3        Econ16.lag3       Econ17.lag3       Econ18.lag3
##          NA            NA            NA            NA
##          Econ19.lag3        Econ1.lag4        Econ2.lag4       Econ3.lag4
##          NA            NA            NA            NA
##          Econ4.lag4        Econ5.lag4        Econ6.lag4       Econ7.lag4
##          NA            NA            NA            NA
##          Econ8.lag4        Econ9.lag4        Econ10.lag4      Econ11.lag4
##          NA            NA            NA            NA
##          Econ12.lag4        Econ13.lag4       Econ14.lag4      Econ15.lag4
##          NA            NA            NA            NA
##          Econ16.lag4        Econ17.lag4       Econ18.lag4      Econ19.lag4
##          NA            NA            NA            NA

```

The dataset includes of 19 economic variables in 5 time lags, where each time lag represents the number of quarters before the start of the specific project. These lagged versions of the same economic variable are highly correlated, and therefore having all lags in the model together leads to multicollinearity. This causes that the model cannot solve for unique coefficients and produced NAs.

c)

```

m1_pred_train <- predict(m1, newdata = df[train,])
rmse_train <- sqrt(mean((df[train, "y"] - m1_pred_train)^2))

```

```
rmse_test <- sqrt(mean((df[test, "y"] - m1_pred_test)^2))
cat('RMSE of training set:', rmse_train, '\n')
```

RMSE of training set: 0.1759113


```
cat('RMSE of test set:', rmse_test)
```

RMSE of test set: 0.3063756

The full model fits the training data closely but performs worse on the test data, which is to be expected. However, the relative increase of $\frac{0.306}{0.176} \approx 1.74$ may be suggesting limited generalization and potential overfitting due to the large number of predictors and multicollinearity.

d)

```
par(mfrow = c(2, 1), cex = 0.7, mar = c(4, 4, 2, 2))
plot(m1_pred_train, df[train, 'y'], ylab = 'Observed', xlab = 'Predicted',
     main = 'Train Data')
abline(0,1,lty=2, col = 'red')
plot(m1_pred_test, df[test, 'y'], ylab = 'Observed', xlab = 'Predicted',
     main = 'Test Data')
abline(0,1,lty=2, col = 'red')
```



Model fits training data relatively well, points lie close to the 45° line. The model, however, shows moderate generalization loss as the test data appears more scattered around the line, which is consistent with earlier RMSE gap. In both plots we can see that at extremes, points shrink toward the line center, indicating slight underprediction of very high values (*Observed > Predicted*) and overprediction of very low values (*Observed < Predicted*).

Exercise 2

Forward Selection

```
m0<- lm(y~1, data = df, subset = train)
m_fs<- step(m0, scope = list(lower = formula(m0), upper = formula(m1)),
            direction = "forward", trace = 0)
summary(m_fs)

##
## Call:
## lm(formula = y ~ PhysFin8 + Econ14.lag4 + PhysFin1 + COMPLETION.YEAR +
##      Econ7.lag3 + PhysFin3 + PhysFin4 + PhysFin6 + PhysFin5 +
##      Econ12.lag4 + COMPLETION.QUARTER, data = df, subset = train)
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -0.60256 -0.11499  0.00195  0.14862  0.56762
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -3.944e+00  1.119e+00 -3.526 0.000507 ***
## PhysFin8     4.798e-04  3.095e-05 15.502 < 2e-16 ***
## Econ14.lag4   6.405e-05  1.779e-05  3.599 0.000389 ***
## PhysFin1     -3.291e-02  3.525e-03 -9.336 < 2e-16 ***
## COMPLETION.YEAR 1.225e-01  1.530e-02  8.002 5.51e-14 ***
## Econ7.lag3    -4.444e-03  1.192e-03 -3.728 0.000242 ***
## PhysFin3     1.438e-04  5.761e-05  2.496 0.013258 *  
## PhysFin4     -5.650e-05  4.553e-05 -1.241 0.215812  
## PhysFin6     5.292e-04  8.879e-05  5.960 9.12e-09 *** 
## PhysFin5     -3.033e-03  5.337e-04 -5.682 3.90e-08 *** 
## Econ12.lag4   3.788e-04  1.145e-04  3.308 0.001087 ** 
## COMPLETION.QUARTER 3.598e-02  1.305e-02  2.757 0.006288 ** 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2216 on 236 degrees of freedom
## Multiple R-squared:  0.9286, Adjusted R-squared:  0.9252 
## F-statistic: 278.9 on 11 and 236 DF,  p-value: < 2.2e-16
```

Backward Selection

```
m_bs<- step(m1, direction = "backward", trace = 0)
summary(m_bs)
```

```

## 
## Call:
## lm(formula = y ~ START.YEAR + COMPLETION.YEAR + COMPLETION.QUARTER +
##     PhysFin1 + PhysFin2 + PhysFin4 + PhysFin5 + PhysFin6 + PhysFin8 +
##     Econ1 + Econ2 + Econ3 + Econ5 + Econ6 + Econ7 + Econ8 + Econ9 +
##     Econ10 + Econ11 + Econ12 + Econ13 + Econ14 + Econ15 + Econ16 +
##     Econ18 + Econ19 + Econ1.lag1 + Econ2.lag1 + Econ3.lag1 +
##     Econ4.lag1 + Econ5.lag1 + Econ6.lag1 + Econ7.lag1 + Econ8.lag1 +
##     Econ10.lag1 + Econ11.lag1 + Econ15.lag1 + Econ16.lag1 + Econ17.lag1 +
##     Econ18.lag1 + Econ19.lag1 + Econ1.lag2 + Econ2.lag2 + Econ3.lag2 +
##     Econ4.lag2 + Econ5.lag2 + Econ6.lag2 + Econ7.lag2 + Econ8.lag2 +
##     Econ9.lag2 + Econ10.lag2 + Econ11.lag2 + Econ12.lag2 + Econ14.lag2 +
##     Econ15.lag2 + Econ16.lag2 + Econ17.lag2 + Econ18.lag2 + Econ19.lag2 +
##     Econ1.lag3 + Econ2.lag3 + Econ3.lag3 + Econ4.lag3, data = df,
##     subset = train)
## 
## Residuals:
##      Min        1Q     Median        3Q       Max
## -0.56748 -0.10845  0.00134  0.10975  0.48532
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)            3.441e+02  8.088e+01  4.254 3.34e-05 ***
## START.YEAR           -5.205e+00  1.230e+00 -4.233 3.64e-05 ***
## COMPLETION.YEAR       7.513e-02  3.053e-02  2.461 0.014764 *  
## COMPLETION.QUARTER   2.626e-02  1.541e-02  1.704 0.090061 .  
## PhysFin1             -2.938e-02  3.690e-03 -7.962 1.70e-13 ***
## PhysFin2              5.576e-05  1.848e-05  3.017 0.002914 ** 
## PhysFin4             -1.270e-04  6.016e-05 -2.111 0.036149 *  
## PhysFin5             -3.016e-03  6.357e-04 -4.744 4.19e-06 ***
## PhysFin6              6.479e-04  1.160e-04  5.585 8.30e-08 *** 
## PhysFin8              4.856e-04  3.137e-05 15.480 < 2e-16 ***
## Econ1                6.132e-04  1.476e-04  4.155 4.98e-05 *** 
## Econ2                5.610e-01  1.255e-01  4.469 1.37e-05 *** 
## Econ3                4.369e-01  9.736e-02  4.488 1.26e-05 *** 
## Econ5                3.399e-05  1.152e-05  2.951 0.003581 ** 
## Econ6                -4.457e-04  2.467e-04 -1.807 0.072392 .  
## Econ7                -1.519e-01  3.778e-02 -4.021 8.44e-05 *** 
## Econ8                1.039e-02  2.917e-03  3.563 0.000468 *** 
## Econ9                -3.002e-04  6.968e-05 -4.308 2.68e-05 *** 
## Econ10               1.005e+00  2.602e-01  3.862 0.000156 *** 
## Econ11               -1.674e-03  4.766e-04 -3.512 0.000560 *** 
## Econ12               3.505e-03  1.141e-03  3.074 0.002437 ** 
## Econ13               -7.891e-05  4.799e-05 -1.644 0.101821 
## Econ14               1.047e-03  2.595e-04  4.033 8.06e-05 *** 
## Econ15               -3.051e-01  9.325e-02 -3.272 0.001276 ** 
## Econ16               7.753e-01  1.961e-01  3.953 0.000110 *** 
## Econ18               2.445e-04  5.446e-05  4.489 1.26e-05 *** 
## Econ19               -4.134e-06  2.468e-06 -1.675 0.095646 .  
## Econ1.lag1            -3.709e-04  1.345e-04 -2.757 0.006424 ** 
## Econ2.lag1            -1.043e+00  2.434e-01 -4.287 2.92e-05 *** 
## Econ3.lag1            -1.384e-01  4.282e-02 -3.233 0.001452 ** 
## Econ4.lag1             4.999e-01  1.863e-01  2.684 0.007945 ** 
## Econ5.lag1             1.306e-05  8.387e-06  1.558 0.121035

```

```

## Econ6.lag1      9.744e-04  2.066e-04   4.717 4.72e-06 ***
## Econ7.lag1     -8.318e-02  1.951e-02  -4.263 3.22e-05 ***
## Econ8.lag1     -4.198e-03  1.394e-03  -3.010 0.002975 **
## Econ10.lag1    -4.502e-01  1.608e-01  -2.799 0.005672 **
## Econ11.lag1    -4.871e-03  1.138e-03  -4.279 3.02e-05 ***
## Econ15.lag1     1.027e+00  2.366e-01   4.340 2.35e-05 ***
## Econ16.lag1    -1.090e+00  2.464e-01  -4.424 1.65e-05 ***
## Econ17.lag1    -5.813e-04  1.917e-04  -3.033 0.002772 **
## Econ18.lag1     1.701e-04  3.887e-05   4.376 2.02e-05 ***
## Econ19.lag1    -9.994e-06  2.355e-06  -4.243 3.49e-05 ***
## Econ1.lag2     -2.018e-04  1.177e-04  -1.715 0.087981 .
## Econ2.lag2      8.928e-01  1.988e-01   4.492 1.24e-05 ***
## Econ3.lag2     -7.571e-02  2.741e-02  -2.763 0.006318 **
## Econ4.lag2     -2.422e-01  1.361e-01  -1.779 0.076847 .
## Econ5.lag2     -3.441e-05  1.667e-05  -2.064 0.040437 *
## Econ6.lag2     -4.755e-04  2.212e-04  -2.150 0.032889 *
## Econ7.lag2     -9.823e-02  2.998e-02  -3.277 0.001256 **
## Econ8.lag2     -1.017e-02  3.573e-03  -2.848 0.004905 **
## Econ9.lag2     -1.561e-04  5.048e-05  -3.093 0.002292 **
## Econ10.lag2    2.591e-01  1.369e-01   1.893 0.059920 .
## Econ11.lag2    -4.081e-03  1.011e-03  -4.035 7.99e-05 ***
## Econ12.lag2    1.061e-02  2.572e-03   4.123 5.66e-05 ***
## Econ14.lag2    -5.948e-04  1.656e-04  -3.593 0.000420 ***
## Econ15.lag2    -2.559e-01  1.102e-01  -2.321 0.021385 *
## Econ16.lag2    8.003e-01  2.489e-01   3.215 0.001539 **
## Econ17.lag2    -3.046e-04  9.671e-05  -3.150 0.001907 **
## Econ18.lag2    2.225e-04  5.154e-05   4.318 2.57e-05 ***
## Econ19.lag2    4.430e-06  1.806e-06   2.453 0.015106 *
## Econ1.lag3     7.778e-04  1.319e-04   5.895 1.76e-08 ***
## Econ2.lag3    -7.609e-01  1.787e-01  -4.257 3.30e-05 ***
## Econ3.lag3     1.231e-01  3.434e-02   3.586 0.000431 ***
## Econ4.lag3    -3.490e-01  1.037e-01  -3.364 0.000933 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2054 on 184 degrees of freedom
## Multiple R-squared:  0.9522, Adjusted R-squared:  0.9358
## F-statistic: 58.14 on 63 and 184 DF,  p-value: < 2.2e-16

```

RSME Comparison

```

pred_test_fs <- predict(m_fs, newdata = df[test,])
pred_test_bs <- predict(m_bs, newdata = df[test,])

rmse_test_fs <- sqrt(mean((df[test,"y"] - pred_test_fs)^2))
rmse_test_bs <- sqrt(mean((df[test,"y"] - pred_test_bs)^2))

cat("RMSE test (forward):", rmse_test_fs, "\n")

## RMSE test (forward): 0.2823906


```

```
cat("RMSE test (backward):", rmse_test_bs)
```

```
## RMSE test (backward): 0.439674
```

Plots Comparison

```
par(mfrow = c(2, 1), cex = 0.7, mar = c(4, 4, 2, 2))
plot(pred_test_fs, df[test,'y'], ylab = 'Observed', xlab = 'Predicted',
      main = 'Forward Selection Model')
abline(0,1,lty=2, col = 'red')
plot(pred_test_bs, df[test,'y'], ylab = 'Observed', xlab = 'Predicted',
      main = 'Backward Selection Model')
abline(0,1,lty=2, col = 'red')
```



The forward selection model has a lower test RMSE (0.28) compared to the backward selection model (0.44), meaning it predicts unseen data more accurately. Compared to the test RMSE of the full model (0.306), the forward selection model shows only slight improvement, while the backward selection model performs worse. In the plots, the forward model's points are closer to the red reference line, while the backward model shows more scatter and a few larger deviations at high predicted values. The difference suggests that backward selection may have retained too many correlated variables, causing higher variance and poorer generalization.

Exercise 3

The forward selection model would be preferred due to its lower test error and simpler structure. ANOVA makes sense only if the models are nested, and as the backward and forward selection models are not nested, we can only compare them individually with the full model. In this way, we can test whether the simpler model explains the data just as well as the full model.

```
#Forward Selection Model vs Full Model Variance Table:  
print(anova(m_fs, m1)[, c("Res.Df", "RSS", "Df", "Sum of Sq", "F", "Pr(>F)")] )
```

```
##   Res.Df     RSS Df Sum of Sq      F Pr(>F)  
## 1    236 11.5905  
## 2    175  7.6743 61    3.9162 1.464 0.02906 *  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#Backward Selection Model vs Full Model Variance Table:  
print(anova(m_bs, m1)[, c("Res.Df", "RSS", "Df", "Sum of Sq", "F", "Pr(>F)")] )
```

```
##   Res.Df     RSS Df Sum of Sq      F Pr(>F)  
## 1    184  7.7613  
## 2    175  7.6743  9   0.086951 0.2203 0.9913
```

The full model explains more variation than the forward model. However, the improvement seems small relative to the complexity added (61 extra variables). The RSS improves for the full model by $11.5905 - 7.6743 \approx 3.92$ that is $\frac{3.9162}{61} = 0.0642$ RSS improvement per added variable. For this, it seems reasonable to continue with the forward selection model. The full model does not improve fit over the backward model. Both perform almost identically in terms of explained variance.

Exercise 4


```
library(cvTools)  
  
cv_m1_train <- cvFit(m1, y = df[train, "y"], data = df[train, ],  
                      cost = rmspe, K = 5, R = 100)  
cv_m1_test <- cvFit(m1, y = df[test, "y"], data = df[test, ],  
                      cost = rmspe, K = 5, R = 100)  
cv_fs_train <- cvFit(m_fs, y = df[train, "y"], data = df[train, ],  
                      cost = rmspe, K = 5, R = 100)  
cv_fs_test <- cvFit(m_fs, y = df[test, "y"], data = df[test, ],  
                      cost = rmspe, K = 5, R = 100)  
cv_bs_train <- cvFit(m_bs, y = df[train, "y"], data = df[train, ],  
                      cost = rmspe, K = 5, R = 100)  
cv_bs_test <- cvFit(m_bs, y = df[test, "y"], data = df[test, ],  
                      cost = rmspe, K = 5, R = 100)
```

```
par(cex = 0.7)  
  
boxplot(list(  
  Full_Train=cv_m1_train$reps, Full_Test=cv_m1_test$reps,
```

```

Forward_Train=cv_fs_train$reps, Forward_Test=cv_fs_test$reps,
Backward_Train=cv_bs_train$reps, Backward_Test=cv_bs_test$reps
), ylab="RMSE", main="5-fold CV (100 reps)")


```



```

boxplot(list(
  Full_Train=cv_m1_train$reps, Full_Test=cv_m1_test$reps,
  Forward_Train=cv_fs_train$reps, Forward_Test=cv_fs_test$reps,
  Backward_Train=cv_bs_train$reps, Backward_Test=cv_bs_test$reps
), ylab="RMSE", main="5-fold CV (100 reps)", ylim = c(0, 40))

```




From the plots it is clear that the forward selection model shows the lowest and most stable RMSE for both training and test sets. Its boxes are compact and centered around the median values, while the full and backward models show extreme variance and instability with higher median RSMEs.

Exercise 5

```
par(cex = 0.7)

cv_m1_train <- cvFit(m1, y = df[train, "y"], data = df[train,],
                      cost = rtmspe, K = 5, R = 100)
cv_m1_test  <- cvFit(m1, y = df[test, "y"], data = df[test, ],
                      cost = rtmspe, K = 5, R = 100)
cv_fs_train <- cvFit(m_fs, y = df[train, "y"], data = df[train,],
                      cost = rtmspe, K = 5, R = 100)
cv_fs_test  <- cvFit(m_fs, y = df[test, "y"], data = df[test, ],
                      cost = rtmspe, K = 5, R = 100)
cv_bs_train <- cvFit(m_bs, y = df[train, "y"], data = df[train,],
                      cost = rtmspe, K = 5, R = 100)
cv_bs_test  <- cvFit(m_bs, y = df[test, "y"], data = df[test, ],
                      cost = rtmspe, K = 5, R = 100)

boxplot(list(
  Full_Train=cv_m1_train$reps, Full_Test=cv_m1_test$reps,
  Forward_Train=cv_fs_train$reps, Forward_Test=cv_fs_test$reps,
  Backward_Train=cv_bs_train$reps, Backward_Test=cv_bs_test$reps
), ylab="RRMSE", main="5-fold CV (100 reps)")
```



Here instead of measuring root mean squared prediction error (RMSE) in absolute units, we measure relative root mean squared prediction error (RRMSE) where each error is scaled by the observed value. That means we observe percentage deviation, which makes comparisons less sensitive to the response's scale. Using

relative RRMSE confirms the same result as before, the forward selection model remains the preferred model.