

Exercise 2

Advanced Methods for Regression and Classification

Rita Selimi

10/28/2024

Load the necessary libraries

Load Data

```
# Load the dataset
load("building.RData")

# Check the data
# head(df)
# Check the structure of the data
str(df)

## 'data.frame': 372 obs. of 108 variables:
## $ y : num 7.7 8.52 7.09 5.11 8.61 ...
## $ START.YEAR : num 81 84 78 72 87 87 87 88 76 80 ...
## $ START.QUARTER : num 1 1 1 2 1 1 2 1 3 1 ...
## $ COMPLETION.YEAR : num 85 89 81 73 90 90 90 89 77 80 ...
## $ COMPLETION.QUARTER: num 1 4 4 2 2 1 1 3 4 4 ...
## $ PhysFin1 : num 1 1 1 1 1 1 1 1 1 ...
## $ PhysFin2 : num 3150 7600 4800 685 3000 2500 1810 1150 2110 3030 ...
## $ PhysFin3 : num 920 1140 840 202 800 640 492 380 540 930 ...
## $ PhysFin4 : num 598.5 3040 480 13.7 1230 ...
## $ PhysFin5 : num 190 400 100 20 410 420 640 500 90 170 ...
## $ PhysFin6 : num 1011 964 690 460 632 ...
## $ PhysFin7 : num 16 23 15 4 13 12 11 6 5 3 ...
## $ PhysFin8 : num 1200 2900 630 140 5000 4800 5700 5300 690 1500 ...
## $ Econ1 : num 6713 3152 1627 2581 6790 ...
## $ Econ2 : num 56.2 106 41 12.1 203.8 ...
## $ Econ3 : num 61.5 103 41.2 10 162.8 ...
## $ Econ4 : num 6.11 3.15 1.74 1.24 6.46 6.46 6.73 3.44 2.28 5.97 ...
## $ Econ5 : num 320957 685698 160402 38194 1640293 ...
## $ Econ6 : num 3486 3526 1218 287 10855 ...
## $ Econ7 : num 64.5 105.5 34.4 13.6 229.3 ...
## $ Econ8 : num 240 209 286 17 393 ...
## $ Econ9 : num 12457 17584 6489 154 69445 ...
## $ Econ10 : int 15 15 15 12 11 11 11 15 15 ...
## $ Econ11 : num 797 1408 614 184 2739 ...
## $ Econ12 : num 810 1474 608 211 3148 ...
```

```

## $ Econ13      : num 1755 8842 1755 1613 9248 ...
## $ Econ14      : num 8003 8864 7773 1649 9380 ...
## $ Econ15      : num 67.8 105.5 45.9 11.6 158.6 ...
## $ Econ16      : num 63.2 105.3 38.3 10.1 169.5 ...
## $ Econ17      : num 3759 12113 1538 393 10082 ...
## $ Econ18      : num 42587 45966 39066 8436 49572 ...
## $ Econ19      : num 628133 1188996 524765 141543 2318397 ...
## $ Econ1.lag1   : num 4986 2700 1580 2952 6370 ...
## $ Econ2.lag1   : num 55.5 103 40.3 11.6 190.3 ...
## $ Econ3.lag1   : num 60.8 101.8 40.8 8.5 154.4 ...
## $ Econ4.lag1   : num 3.94 2.65 1.15 1.99 5.33 5.33 6.46 3.8 2.32 4.3 ...
## $ Econ5.lag1   : num 297210 625829 150267 35859 1523167 ...
## $ Econ6.lag1   : num 3664 4387 1150 322 12930 ...
## $ Econ7.lag1   : num 61.5 100.4 34.1 12.7 210.7 ...
## $ Econ8.lag1   : num 179.6 156.6 214.3 56.6 295 ...
## $ Econ9.lag1   : num 9342 13188 4867 610 52084 ...
## $ Econ10.lag1  : int 15 15 15 12 11 11 11 11 15 15 ...
## $ Econ11.lag1  : num 758 1424 574 165 2595 ...
## $ Econ12.lag1  : num 862 1584 680 209 3000 ...
## $ Econ13.lag1  : num 1755 8777 1755 1504 9330 ...
## $ Econ14.lag1  : num 8018 8799 6714 1582 9396 ...
## $ Econ15.lag1  : num 65 101 43.4 10.9 148.8 ...
## $ Econ16.lag1  : num 60.53 101.89 36.45 9.79 159 ...
## $ Econ17.lag1  : num 3539 13572 1535 435 9700 ...
## $ Econ18.lag1  : num 31940 34475 29300 32776 37179 ...
## $ Econ19.lag1  : num 610503 1067772 466212 129102 1908976 ...
## $ Econ1.lag2   : num 6788 3561 2628 2649 5909 ...
## $ Econ2.lag2   : num 54.2 98.2 39.3 11.4 177.6 ...
## $ Econ3.lag2   : num 59.4 98.64 40.21 6.97 147.44 ...
## $ Econ4.lag2   : num 5.41 2.76 1.52 2.25 6.88 6.88 5.33 6.54 2.69 3.53 ...
## $ Econ5.lag2   : num 280452 602225 143738 32794 1451176 ...
## $ Econ6.lag2   : num 3756 3819 1284 389 8146 ...
## $ Econ7.lag2   : num 58.1 97.2 33.5 11.7 188.9 ...
## $ Econ8.lag2   : num 119.8 104.4 142.9 42.5 196.7 ...
## $ Econ9.lag2   : num 6228 8792 3245 458 34722 ...
## $ Econ10.lag2  : int 15 15 15 12 11 11 11 11 15 15 ...
## $ Econ11.lag2  : num 795 1299 554 168 2284 ...
## $ Econ12.lag2  : num 818 1390 664 210 2628 ...
## $ Econ13.lag2  : num 1755 8700 1755 1450 9297 ...
## $ Econ14.lag2  : num 8001 8735 5827 1507 9347 ...
## $ Econ15.lag2  : num 63.7 98.1 41.8 10.2 140.9 ...
## $ Econ16.lag2  : num 58.55 98.45 34.76 9.35 146.2 ...
## $ Econ17.lag2  : num 3348 13596 1528 509 10149 ...
## $ Econ18.lag2  : num 21294 22983 19533 24582 24786 ...
## $ Econ19.lag2  : num 589390 973524 409678 123618 1681849 ...
## $ Econ1.lag3   : num 5728 3157 2374 2312 7045 ...
## $ Econ2.lag3   : num 52.4 92.8 38 10.6 160 ...
## $ Econ3.lag3   : num 57.65 96.49 39.43 5.44 141.34 ...
## $ Econ4.lag3   : num 5.4 3.05 0.92 2.58 4.72 4.72 6.88 6.73 2.68 3.39 ...
## $ Econ5.lag3   : num 262789 552124 134548 30012 1341073 ...
## $ Econ6.lag3   : num 2931 3897 1191 345 8245 ...
## $ Econ7.lag3   : num 54.2 96.9 33.7 10.8 173.8 ...
## $ Econ8.lag3   : num 59.9 52.2 71.5 28.3 98.3 ...
## $ Econ9.lag3   : num 3114 4396 1622 305 17361 ...

```

```

## $ Econ10.lag3      : int 15 15 15 12 11 11 11 11 15 15 ...
## $ Econ11.lag3     : num 747 1294 575 180 2451 ...
## $ Econ12.lag3     : num 816 1288 680 158 2526 ...
## $ Econ13.lag3     : num 1755 8556 1755 1439 9254 ...
## $ Econ14.lag3     : num 8013 8585 5565 1450 9306 ...
## $ Econ15.lag3     : num 62.78 95.35 41.03 9.91 136.56 ...
## $ Econ16.lag3     : num 56.45 94.34 33.37 8.85 138.8 ...
## $ Econ17.lag3     : num 3388 12064 1602 591 9291 ...
## $ Econ18.lag3     : num 10647 11492 9766 16388 12393 ...
## $ Econ19.lag3     : num 606524 954629 403875 121857 1732938 ...
## $ Econ1.lag4      : num 7196 3678 2693 1381 5606 ...
## $ Econ2.lag4      : num 51.3 86.2 36.2 10 149.1 ...
## $ Econ3.lag4      : num 56.13 83.21 37.64 3.91 134.8 ...
## $ Econ4.lag4      : num 5.97 3.25 1.55 3 4.09 4.09 4.72 6.46 3.56 3.25 ...
## $ Econ5.lag4      : num 249111 526596 134313 27231 1284199 ...
## $ Econ6.lag4      : num 2562 2791 1529 316 6622 ...
## $ Econ7.lag4      : num 52.8 94.1 31.43 9.85 147.6 ...
## $ Econ8.lag4      : num 217 334.8 175.7 14.2 432.4 ...
## $ Econ9.lag4      : num 10446 14489 3995 153 73144 ...
## $ Econ10.lag4     : int 15 15 15 12 14 14 11 11 15 15 ...
## [list output truncated]

```

Data Splitting

```

set.seed(12332281)

# Split the data
sample_index <- sample(1:nrow(df), size = 2/3 * nrow(df))
train_data <- df[sample_index, ]
test_data <- df[-sample_index, ]

# Display the number of samples in each set
cat("Training set size:", nrow(train_data), "\n")

## Training set size: 248

cat("Test set size:", nrow(test_data), "\n")

## Test set size: 124

```

1. Fit regression model on the training data

```

# Fit linear regression model
model <- lm(y ~ ., data = train_data)
summary(model)

##
## Call:

```

```

## lm(formula = y ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -0.62167 -0.10676  0.01542  0.13256  0.79241 
## 
## Coefficients: (35 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 5.462e+01 2.013e+02  0.271  0.7864  
## START.YEAR -6.769e-01 3.259e+00 -0.208  0.8357  
## START.QUARTER -1.424e+00 2.867e+00 -0.497  0.6201  
## COMPLETION.YEAR 9.123e-02 3.554e-02  2.567  0.0111 *  
## COMPLETION.QUARTER 4.704e-02 1.859e-02  2.531  0.0123 *  
## PhysFin1     -2.910e-02 4.714e-03 -6.173 4.53e-09 *** 
## PhysFin2      1.068e-04 4.881e-05  2.187  0.0301 *  
## PhysFin3     -3.177e-04 1.390e-04 -2.285  0.0235 *  
## PhysFin4     -9.701e-05 6.768e-05 -1.433  0.1535  
## PhysFin5     -3.698e-03 6.892e-04 -5.365 2.54e-07 *** 
## PhysFin6      7.609e-04 1.323e-04  5.752 3.86e-08 *** 
## PhysFin7          NA        NA        NA        NA      
## PhysFin8      5.139e-04 3.545e-05 14.496 < 2e-16 *** 
## Econ1         8.098e-06 2.736e-04  0.030  0.9764  
## Econ2        -2.633e-01 7.276e-01 -0.362  0.7178  
## Econ3         4.186e-01 1.906e-01  2.196  0.0294 *  
## Econ4         2.023e-01 2.335e-01  0.867  0.3873  
## Econ5        -3.151e-05 5.995e-05 -0.526  0.5999  
## Econ6        -5.921e-04 5.914e-04 -1.001  0.3181  
## Econ7        -1.800e-01 8.049e-02 -2.237  0.0266 *  
## Econ8         1.543e-02 9.111e-03  1.694  0.0921 .  
## Econ9        -2.704e-04 1.256e-04 -2.153  0.0327 *  
## Econ10        3.662e-01 4.471e-01  0.819  0.4139  
## Econ11        -2.477e-03 1.112e-03 -2.226  0.0273 *  
## Econ12        1.944e-03 3.712e-03  0.524  0.6011  
## Econ13        -3.598e-04 2.656e-04 -1.355  0.1772  
## Econ14        1.539e-04 9.799e-04  0.157  0.8754  
## Econ15        -1.307e-01 1.246e-01 -1.049  0.2956  
## Econ16        1.104e+00 6.323e-01  1.746  0.0826 .  
## Econ17        1.996e-04 2.960e-04  0.674  0.5010  
## Econ18        -8.322e-05 4.037e-04 -0.206  0.8369  
## Econ19        -1.530e-05 8.141e-06 -1.880  0.0618 .  
## Econ1.lag1    -8.527e-04 5.212e-04 -1.636  0.1036  
## Econ2.lag1    -6.409e-01 6.169e-01 -1.039  0.3003  
## Econ3.lag1    -1.559e-01 9.343e-02 -1.668  0.0970 .  
## Econ4.lag1    1.078e+00 5.670e-01  1.901  0.0589 .  
## Econ5.lag1    1.898e-05 2.076e-05  0.914  0.3619  
## Econ6.lag1    6.046e-04 3.616e-04  1.672  0.0963 .  
## Econ7.lag1    -5.435e-02 5.815e-02 -0.935  0.3512  
## Econ8.lag1    -1.347e-02 6.783e-03 -1.986  0.0486 *  
## Econ9.lag1    1.410e-04 1.682e-04  0.838  0.4031  
## Econ10.lag1   -5.527e-01 3.228e-01 -1.712  0.0886 .  
## Econ11.lag1   -2.559e-03 2.798e-03 -0.914  0.3617  
## Econ12.lag1   -5.455e-04 2.616e-03 -0.209  0.8351  
## Econ13.lag1   3.317e-04 3.851e-04  0.861  0.3903  
## Econ14.lag1   1.026e-03 1.467e-03  0.699  0.4852

```

## Econ15.lag1	5.226e-01	4.303e-01	1.214	0.2262
## Econ16.lag1	-8.174e-01	4.237e-01	-1.929	0.0553 .
## Econ17.lag1	-5.078e-04	3.945e-04	-1.287	0.1997
## Econ18.lag1	-1.115e-04	3.938e-04	-0.283	0.7773
## Econ19.lag1	-8.797e-06	6.767e-06	-1.300	0.1953
## Econ1.lag2	-5.505e-04	2.759e-04	-1.995	0.0475 *
## Econ2.lag2	1.073e+00	6.536e-01	1.642	0.1024
## Econ3.lag2	-8.127e-02	9.709e-02	-0.837	0.4037
## Econ4.lag2	-3.551e-01	7.570e-01	-0.469	0.6396
## Econ5.lag2	7.792e-05	9.835e-05	0.792	0.4293
## Econ6.lag2	5.508e-05	4.930e-04	0.112	0.9112
## Econ7.lag2	8.836e-03	7.608e-02	0.116	0.9077
## Econ8.lag2	-1.685e-02	7.866e-03	-2.142	0.0336 *
## Econ9.lag2	4.566e-05	2.446e-04	0.187	0.8522
## Econ10.lag2	5.226e-01	3.850e-01	1.357	0.1764
## Econ11.lag2	-2.943e-03	1.547e-03	-1.902	0.0588 .
## Econ12.lag2	3.544e-03	6.200e-03	0.572	0.5683
## Econ13.lag2	1.484e-04	1.965e-04	0.755	0.4510
## Econ14.lag2	-4.028e-04	3.317e-04	-1.214	0.2263
## Econ15.lag2	-3.784e-01	3.156e-01	-1.199	0.2322
## Econ16.lag2	-2.168e-01	7.170e-01	-0.302	0.7628
## Econ17.lag2	2.241e-04	4.399e-04	0.509	0.6111
## Econ18.lag2	-4.286e-05	3.691e-04	-0.116	0.9077
## Econ19.lag2	6.054e-06	4.401e-06	1.375	0.1708
## Econ1.lag3	7.235e-04	3.441e-04	2.102	0.0370 *
## Econ2.lag3	-4.822e-01	2.840e-01	-1.698	0.0913 .
## Econ3.lag3	1.581e-01	1.368e-01	1.155	0.2495
## Econ4.lag3	-2.271e-01	2.762e-01	-0.822	0.4120
## Econ5.lag3	NA	NA	NA	NA
## Econ6.lag3	NA	NA	NA	NA
## Econ7.lag3	NA	NA	NA	NA
## Econ8.lag3	NA	NA	NA	NA
## Econ9.lag3	NA	NA	NA	NA
## Econ10.lag3	NA	NA	NA	NA
## Econ11.lag3	NA	NA	NA	NA
## Econ12.lag3	NA	NA	NA	NA
## Econ13.lag3	NA	NA	NA	NA
## Econ14.lag3	NA	NA	NA	NA
## Econ15.lag3	NA	NA	NA	NA
## Econ16.lag3	NA	NA	NA	NA
## Econ17.lag3	NA	NA	NA	NA
## Econ18.lag3	NA	NA	NA	NA
## Econ19.lag3	NA	NA	NA	NA
## Econ1.lag4	NA	NA	NA	NA
## Econ2.lag4	NA	NA	NA	NA
## Econ3.lag4	NA	NA	NA	NA
## Econ4.lag4	NA	NA	NA	NA
## Econ5.lag4	NA	NA	NA	NA
## Econ6.lag4	NA	NA	NA	NA
## Econ7.lag4	NA	NA	NA	NA
## Econ8.lag4	NA	NA	NA	NA
## Econ9.lag4	NA	NA	NA	NA
## Econ10.lag4	NA	NA	NA	NA
## Econ11.lag4	NA	NA	NA	NA

```

## Econ12.lag4      NA      NA      NA
## Econ13.lag4      NA      NA      NA
## Econ14.lag4      NA      NA      NA
## Econ15.lag4      NA      NA      NA
## Econ16.lag4      NA      NA      NA
## Econ17.lag4      NA      NA      NA
## Econ18.lag4      NA      NA      NA
## Econ19.lag4      NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2296 on 175 degrees of freedom
## Multiple R-squared:  0.9472, Adjusted R-squared:  0.9255
## F-statistic: 43.61 on 72 and 175 DF,  p-value: < 2.2e-16

```

In the model output, some variables show NA values for their estimates because they are too similar to other variables. This issue, called multicollinearity, happens when certain variables contain almost the same information as others, making it hard for the model to tell their effects apart. To handle this, R leaves out these extra variables by marking them as NA.

a) Comparing Observed vs Predicted Values (training data)

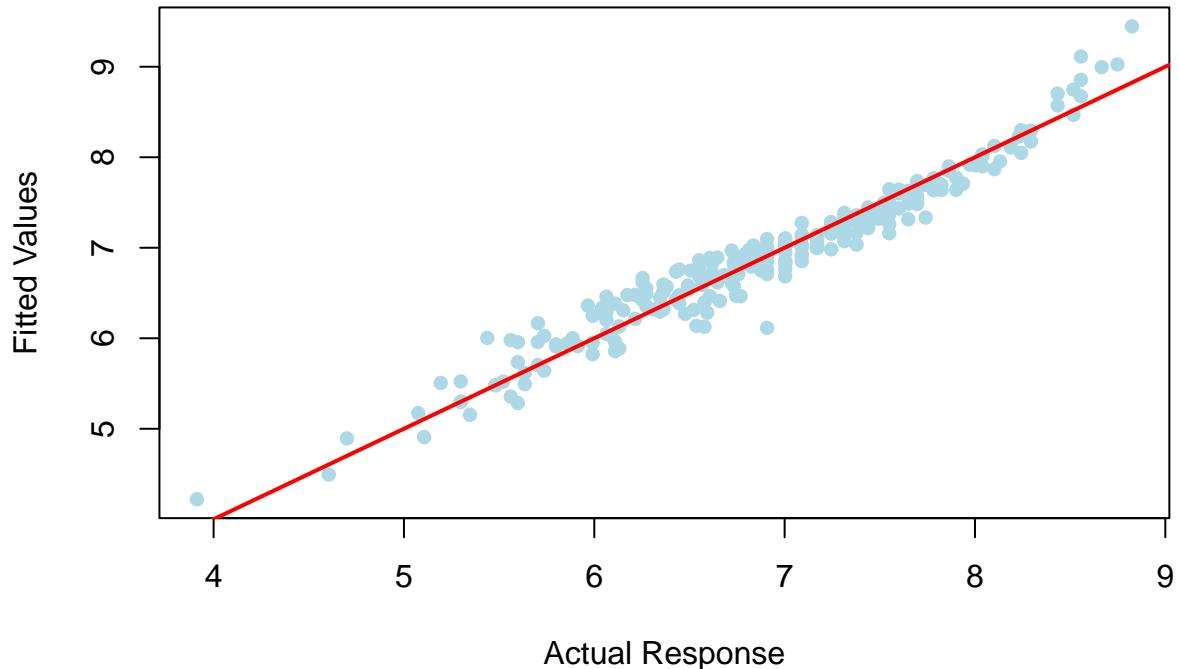
```

train_pred <- predict(model, newdata = train_data)

# Plot for Training Data
plot(train_data$y, train_pred,
      main = "Fitted vs Actual Response",
      xlab = "Actual Response",
      ylab = "Fitted Values",
      pch = 16, col = "lightblue")
abline(0, 1, col = "red", lwd = 2)

```

Fitted vs Actual Response



```
# Compute RMSE for training data
train_rmse <- sqrt(mean((train_data$y - train_pred)^2))
cat("Training RMSE:", train_rmse, "\n")
```

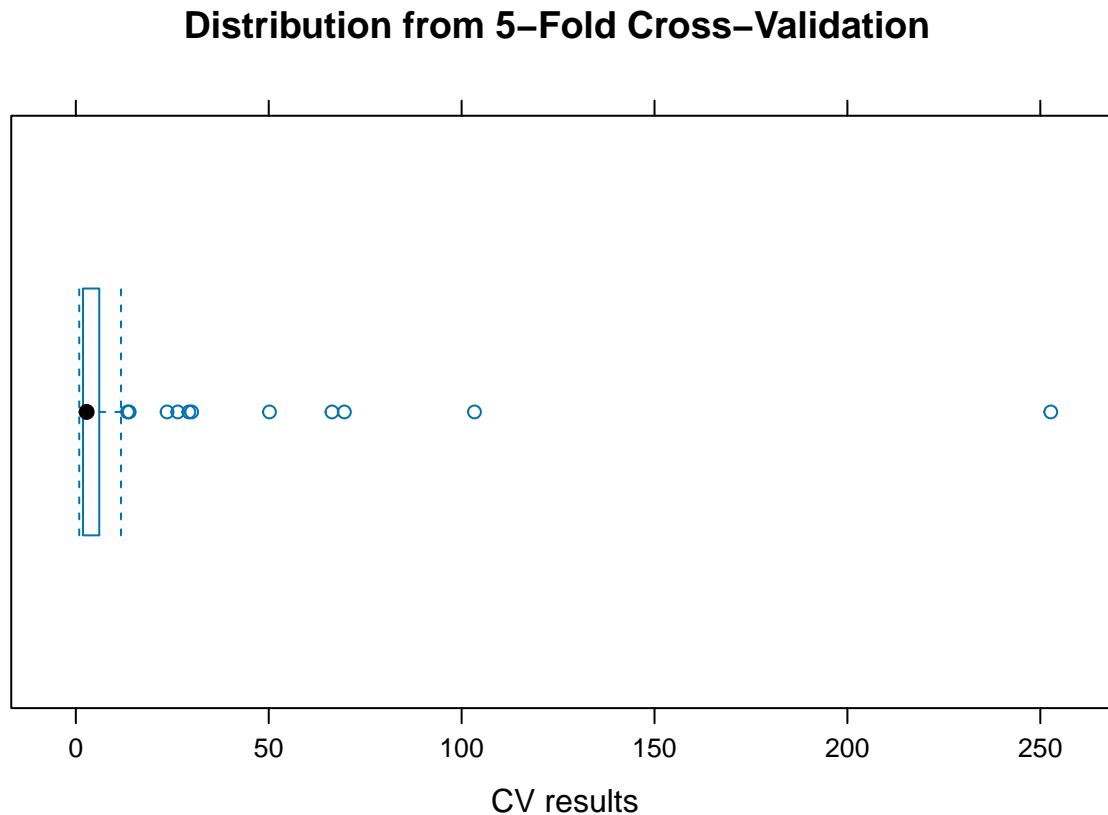
```
## Training RMSE: 0.1928942
```

b) 5-Fold Cross-Validation with 100 Replications (cost=rmspe)

```
# Fit the model using cross-validation
cv_results <- cvFit(
  model,
  data = train_data,
  y = train_data$y,
  cost = rmspe,
  K = 5,
  R = 100
)
print(cv_results)
```

```
## 5-fold CV results:
##      CV
## 9.960448
```

```
plot(cv_results, main = "Distribution from 5-Fold Cross-Validation")
```



The cross-validation plot shows that most of the RMSE values are close to 10, meaning the model generally predicts well across different parts of the data. However, there are a few outliers with much higher error values, suggesting that the model struggles on certain data subsets.

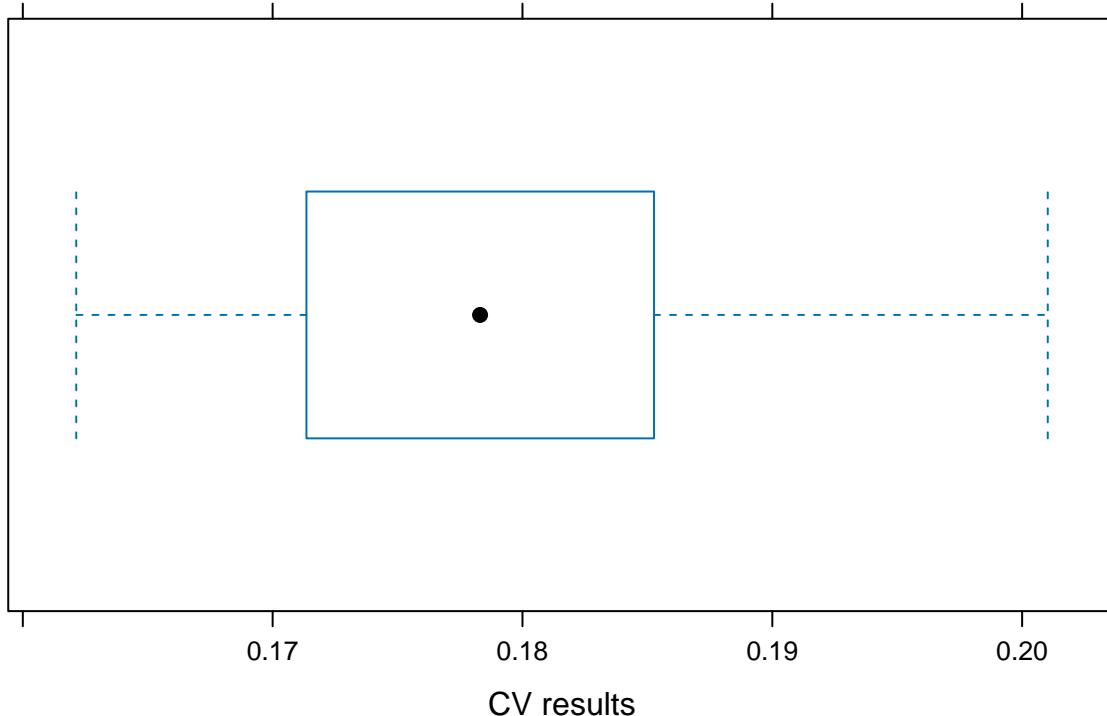
c) 5-Fold Cross-Validation with 100 Replications (cost=rtmspe)

```
# Fit the model using cross-validation
cv_results2 <- cvFit(
  model,
  data = train_data,
  y = train_data$y,
  cost = rtmspe,
  K = 5, # 5-fold cross-validation
  R = 100 # 100 replications
)
print(cv_results2)

## 5-fold CV results:
##      CV
## 0.1790902
```

```
plot(cv_results2, main = "Distribution from 5-Fold Cross-Validation")
```

Distribution from 5-Fold Cross-Validation



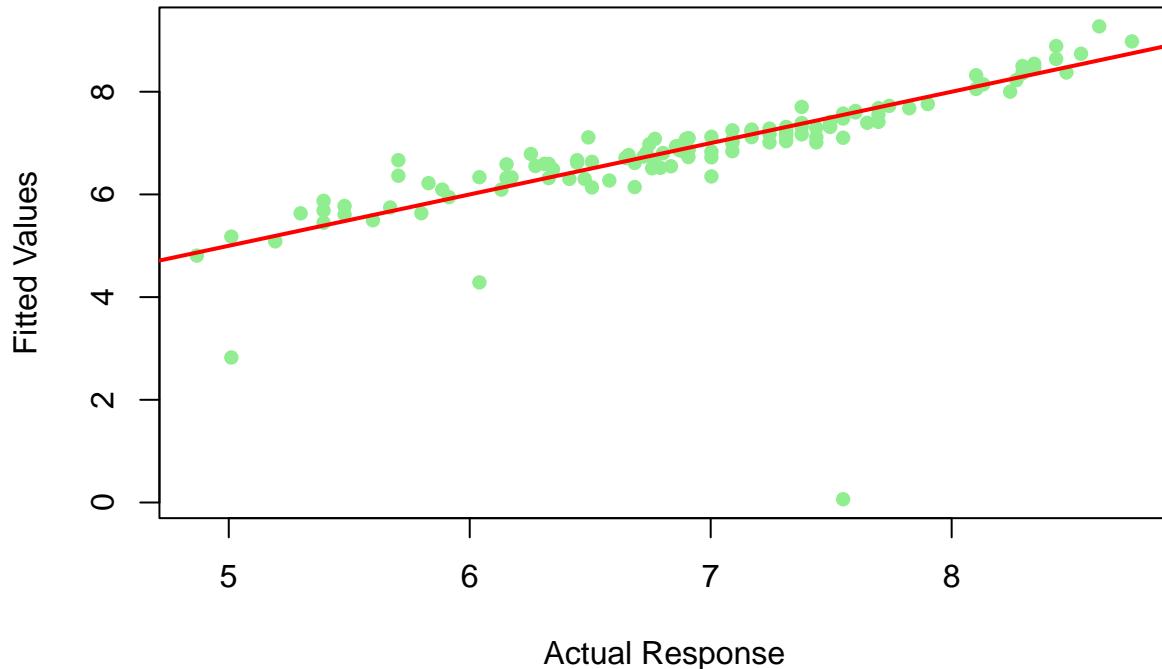
We are using `rtmspe`, an error metric that reduces the influence of outliers, meaning that the focus is on predicting errors and ignores extreme cases that could skew the results. The plot shows that most of the errors fall within a tight range (between 0.17 and 0.20), indicating that the model performs consistently well across different test sets when outliers are down-weighted.

d) Comparing Observed vs Predicted Values (test data)

```
test_pred <- predict(model, newdata = test_data)

# Plot for Test Data
plot(test_data$y, test_pred,
     main = "Fitted vs Actual Response",
     xlab = "Actual Response",
     ylab = "Fitted Values",
     pch = 16, col = "lightgreen")
abline(0, 1, col = "red", lwd = 2)
```

Fitted vs Actual Response



```
# Compute RMSE for testing data
test_rmse <- sqrt(mean((test_data$y - test_pred)^2))
cat("Testing RMSE:", test_rmse, "\n")
```

```
## Testing RMSE: 0.7586301
```

The testing RMSE (0.7586) is higher than the training RMSE (0.1929), suggesting that the model performs well on the training data but may be overfitting, as it is less accurate on new test data. The scatter plot for the test data shows that while most predictions follow the red line (indicating a good fit), there is some spread and a few deviations, which contribute to the higher RMSE.

2. Best Subset Regression

a) Principal Component Regression (PCR)

```
# Reduce predictors to top 50 based on variance
variances <- apply(train_data[, -ncol(train_data)], 2, var)
top_50_vars <- names(sort(variances, decreasing = TRUE))[1:50]
train_data_reduced <- train_data[, c(top_50_vars, "y")]

# Fit PCR model on the reduced set of 50 predictors
pqr_model <- pqr(y ~ ., data = train_data_reduced, scale = TRUE, validation = "CV")
summary(pqr_model)
```

```

## Data: X dimension: 248 50
## Y dimension: 248 1
## Fit method: svdpc
## Number of components considered: 50
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
## (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV 0.8429 0.5917 0.5932 0.5908 0.5422 0.5442 0.5442
## adjCV 0.8429 0.5915 0.5930 0.5901 0.5416 0.5436 0.5440
## 7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV 0.5105 0.5146 0.5225 0.3437 0.3325 0.3172 0.3193
## adjCV 0.5088 0.5131 0.5218 0.3416 0.3307 0.3153 0.3182
## 14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV 0.3187 0.3192 0.3229 0.3233 0.3219 0.3242 0.3250
## adjCV 0.3177 0.3182 0.3218 0.3222 0.3206 0.3229 0.3236
## 21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV 0.3278 0.3246 0.3252 0.3255 0.3265 0.3240 0.3216
## adjCV 0.3263 0.3229 0.3236 0.3239 0.3249 0.3219 0.3196
## 28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV 0.3247 0.3268 0.3314 0.3326 0.3359 0.3323 0.3333
## adjCV 0.3228 0.3249 0.3292 0.3306 0.3334 0.3294 0.3302
## 35 comps 36 comps 37 comps 38 comps 39 comps 40 comps 41 comps
## CV 0.3332 0.3338 0.3343 0.3360 0.3362 0.3401 0.3412
## adjCV 0.3303 0.3310 0.3313 0.3328 0.3331 0.3369 0.3378
## 42 comps 43 comps 44 comps 45 comps 46 comps 47 comps 48 comps
## CV 0.3450 0.3488 0.3499 0.3469 0.3481 0.3520 0.4033
## adjCV 0.3414 0.3447 0.3457 0.3429 0.3440 0.3477 0.3949
## 49 comps 50 comps
## CV 0.4093 0.3927
## adjCV 0.4004 0.3850
##
## TRAINING: % variance explained
## 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X 62.80 71.77 77.59 83.03 86.65 89.14 91.29 93.26
## y 50.82 51.02 53.05 60.54 60.57 60.81 66.68 67.39
## 9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X 94.93 96.06 96.85 97.46 97.96 98.36 98.68
## y 67.53 84.98 85.89 86.99 87.00 87.05 87.10
## 16 comps 17 comps 18 comps 19 comps 20 comps 21 comps 22 comps
## X 98.94 99.11 99.27 99.39 99.49 99.58 99.63
## y 87.10 87.15 87.32 87.32 87.34 87.34 87.58
## 23 comps 24 comps 25 comps 26 comps 27 comps 28 comps 29 comps
## X 99.69 99.73 99.78 99.81 99.84 99.87 99.89
## y 87.60 87.63 87.71 88.03 88.14 88.27 88.29
## 30 comps 31 comps 32 comps 33 comps 34 comps 35 comps 36 comps
## X 99.91 99.93 99.94 99.95 99.96 99.97 99.97
## y 88.30 88.35 88.48 88.89 88.91 88.94 88.95
## 37 comps 38 comps 39 comps 40 comps 41 comps 42 comps 43 comps
## X 99.98 99.99 99.99 99.99 100.00 100.00 100.00
## y 89.07 89.11 89.11 89.14 89.25 89.26 89.36
## 44 comps 45 comps 46 comps 47 comps 48 comps 49 comps 50 comps
## X 100.0 100.00 100.00 100.00 100.00 100.00 100.00
## y 89.4 89.53 89.53 89.53 89.53 89.62 89.74

```

```
# Extract the top 10 components from the PCR model
pcr_scores <- as.data.frame(pcr_model$scores[, 1:10])
pcr_scores$y <- train_data$y # Add the response variable
```

First, I calculated the variance of each predictor in the training data and selected the top 50 predictors with the highest variance. This approach ensures that we retain the predictors most likely to contain significant information. Next, I applied Principal Component Regression (PCR) on this reduced dataset of 50 predictors, which creates new components that capture the main patterns in the data. Finally, we extracted the top 10 components from the PCR model, allowing us to limit the model size to 10 regressors. These components, along with the response variable, form a new dataset that can be used for further modeling.

Best Subset Regression

```
# Run best subset regression with the modified data
best_subset <- regsubsets(y ~ ., data = pcr_scores, nvmax = 10, really.big = TRUE)
summary(best_subset)
```

```
## Subset selection object
## Call: regsubsets.formula(y ~ ., data = pcr_scores, nvmax = 10, really.big = TRUE)
## 10 Variables (and intercept)
##          Forced in Forced out
## `Comp 1`    FALSE    FALSE
## `Comp 2`    FALSE    FALSE
## `Comp 3`    FALSE    FALSE
## `Comp 4`    FALSE    FALSE
## `Comp 5`    FALSE    FALSE
## `Comp 6`    FALSE    FALSE
## `Comp 7`    FALSE    FALSE
## `Comp 8`    FALSE    FALSE
## `Comp 9`    FALSE    FALSE
## `Comp 10`   FALSE   FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##          `Comp 1` `Comp 2` `Comp 3` `Comp 4` `Comp 5` `Comp 6` `Comp 7` 
## 1  ( 1 )    *      " "     " "     " "     " "     " "     " "
## 2  ( 1 )    *      " "     " "     " "     " "     " "     " "
## 3  ( 1 )    *      " "     " "     *      " "     " "     " "
## 4  ( 1 )    *      " "     " "     *      " "     " "     * "
## 5  ( 1 )    *      " "     *      *      " "     " "     * "
## 6  ( 1 )    *      " "     *      *      " "     " "     * "
## 7  ( 1 )    *      " "     *      *      " "     *      * "
## 8  ( 1 )    *      *      *      *      " "     *      * "
## 9  ( 1 )    *      *      *      *      " "     *      * "
## 10 ( 1 )   *      *      *      *      *      " "     * "
##          `Comp 8` `Comp 9` `Comp 10` 
## 1  ( 1 )    " "     " "     " "
## 2  ( 1 )    " "     " "     * "
## 3  ( 1 )    " "     " "     * "
## 4  ( 1 )    " "     " "     * "
## 5  ( 1 )    " "     " "     * "
## 6  ( 1 )    *      " "     * "
```

```

## 7  ( 1 )  "*"      " "
## 8  ( 1 )  "*"      " "
## 9  ( 1 )  "*"      " "
## 10 ( 1 )  "*"      " "

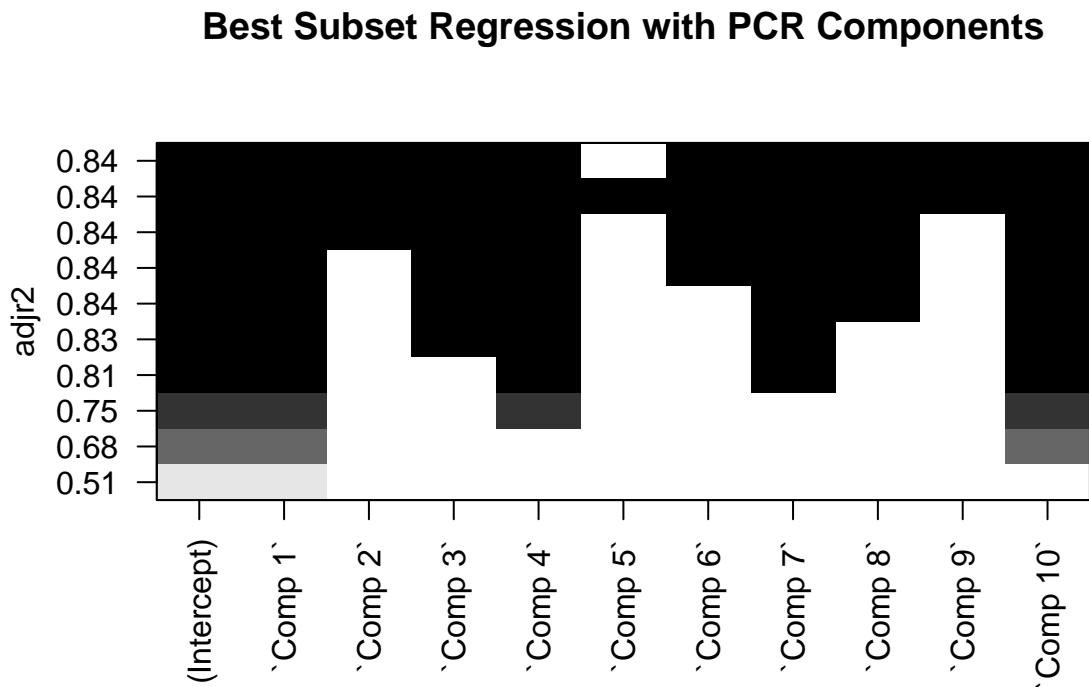
```

b) Plot the Best Subset Results

```

# Plot the best subset regression results to determine the best model
plot(best_subset, scale = "adjr2", main = "Best Subset Regression with PCR Components")

```



In this best subset regression plot, we are comparing models with different combinations of components to find the one with the highest adjusted R^2 , which indicates the model's accuracy while adjusting for the number of components used. The model that seems to be the best includes Comp 1+Comp 3+Comp 4+Comp 7+Comp 8+Comp 10. This model is likely the best choice, as it explains the most variability in the data with a minimal number of components.

c) Apply lm() on the final best model

```

# Fit the linear model using the selected components
final_model <- lm(y ~ `Comp 1` + `Comp 3` + `Comp 4` + `Comp 7` +
                    `Comp 8` + `Comp 10`, data = pcr_scores)
summary(final_model)

```

```

## 
## Call:
## lm(formula = y ~ `Comp 1` + `Comp 3` + `Comp 4` + `Comp 7` +
##     `Comp 8` + `Comp 10`, data = pcr_scores)
## 
## Residuals:
##      Min      1Q Median      3Q      Max 
## -1.36875 -0.16630  0.03118  0.21907  0.80714 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 6.886883  0.021379 322.128 < 2e-16 ***
## `Comp 1`    0.107015  0.003823  27.992 < 2e-16 ***
## `Comp 3`   -0.070281  0.012553 -5.599 5.87e-08 ***
## `Comp 4`    0.139650  0.012994 10.747 < 2e-16 ***
## `Comp 7`    0.196474  0.020663  9.509 < 2e-16 ***
## `Comp 8`    0.071575  0.021596  3.314  0.00106 ** 
## `Comp 10`   -0.466545  0.028437 -16.406 < 2e-16 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 0.3367 on 241 degrees of freedom 
## Multiple R-squared:  0.8437, Adjusted R-squared:  0.8398 
## F-statistic: 216.8 on 6 and 241 DF,  p-value: < 2.2e-16

```

```

# Evaluate the model with 5-fold cross-validation and 100 replications
# (1b) Using rmspe as cost
cv_results_rmspe <- cvFit(
  final_model,
  data = pcr_scores,
  y = pcr_scores$y,
  cost = rmspe,
  K = 5,
  R = 100
)

# Display and plot the cross-validation results for rmspe
print(cv_results_rmspe)

```

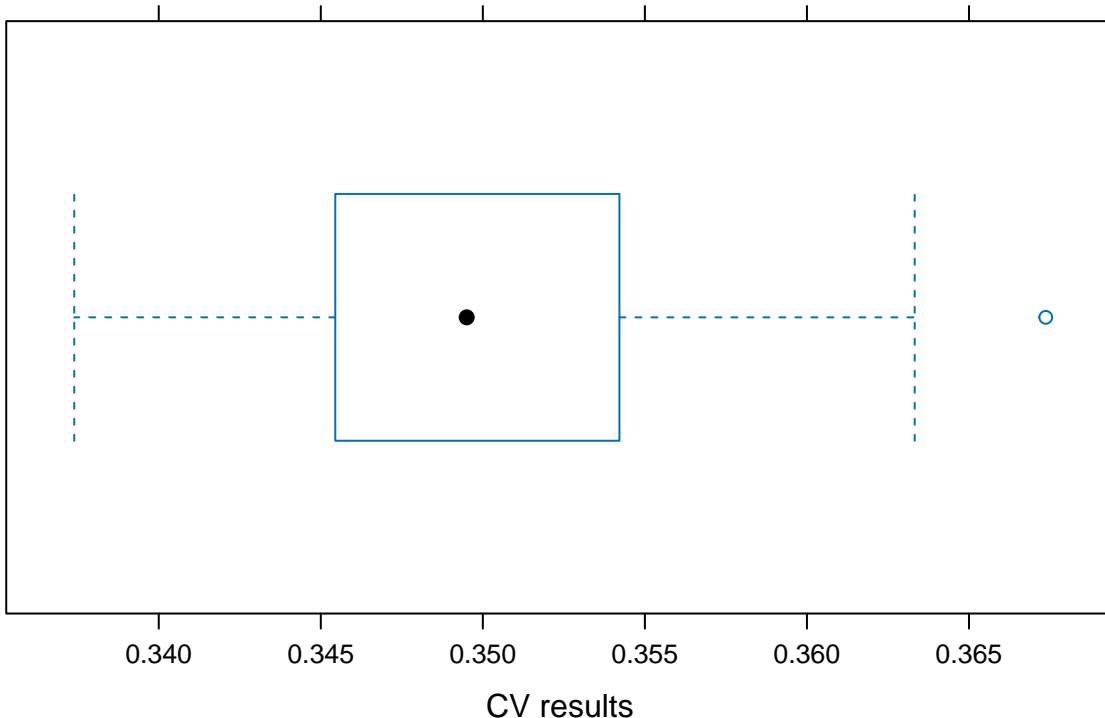
```

## 5-fold CV results:
##      CV
## 0.3501755

plot(cv_results_rmspe, main = "5-Fold Cross-Validation with RMSPE")

```

5–Fold Cross–Validation with RMSPE



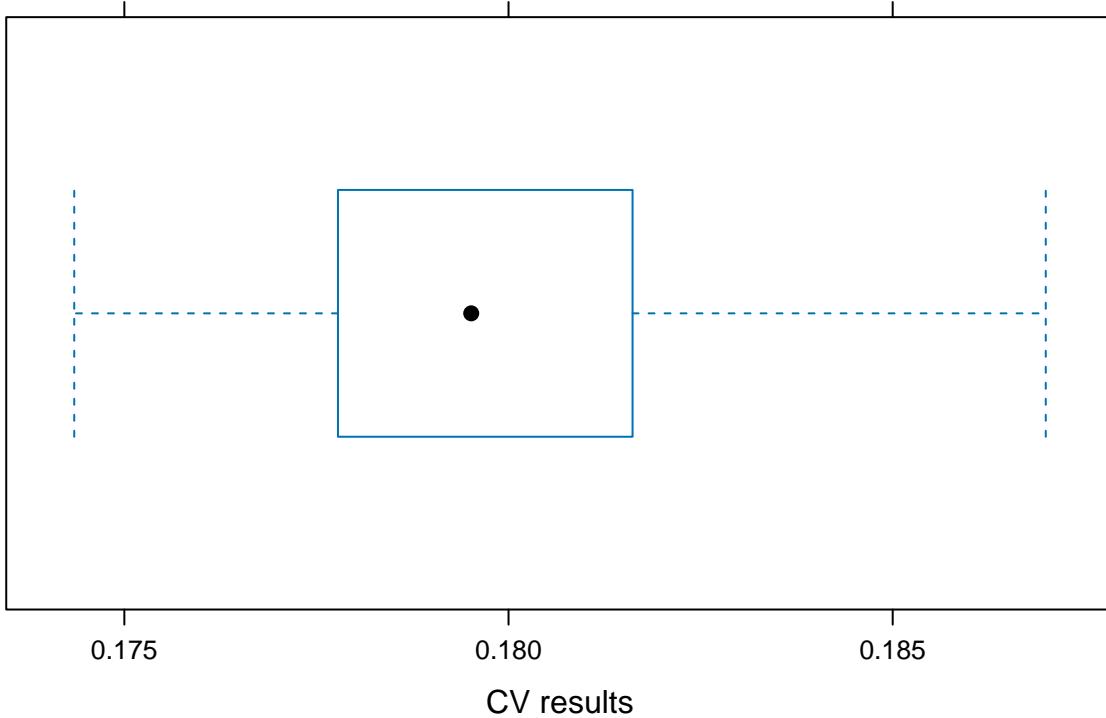
```
# (1c) Using rtmspe as cost
cv_results_rtmspe <- cvFit(
  final_model,
  data = pcr_scores,
  y = pcr_scores$y,
  cost = rtmspe,
  K = 5,
  R = 100
)

# Display and plot the cross-validation results for rtmspe
print(cv_results_rtmspe)

## 5-fold CV results:
##          CV
## 0.1798701

plot(cv_results_rtmspe, main = "5–Fold Cross–Validation with RTMSPE")
```

5-Fold Cross-Validation with RTMSPE



The cross-validation results using two error metrics show that, with RMSPE, the model achieved a cross-validated error of approximately 0.35, and with RTMSPE, the model had a lower error of approximately 0.17. These results indicate that the selected model performs well with a relatively low prediction error.

d) Comparing Observed vs Predicted Values (test data)

```
# Ensure `test_data_reduced` matches `train_data_reduced` in terms of columns and order
test_data_reduced <- test_data[, colnames(train_data_reduced)[-ncol(train_data_reduced)]]

# Manually extract means and sds from `train_data_reduced` if needed
train_means <- apply(train_data_reduced[, -ncol(train_data_reduced)], 2, mean)
train_sds <- apply(train_data_reduced[, -ncol(train_data_reduced)], 2, sd)

# Center and scale `test_data_reduced` using these means and sds
test_data_scaled <- scale(test_data_reduced, center = train_means, scale = train_sds)

# Calculate the components manually by projecting on `pcr_model$loadings`
pcr_test_scores <- as.data.frame(test_data_scaled %*% pcr_model$loadings[, 1:10])

# Set column names to match `pcr_scores`
colnames(pcr_test_scores) <- colnames(pcr_scores)[1:10]

# Ensure pcr_test_scores includes the actual `y` values
pcr_test_scores$y <- test_data$y
```

```

# Predict response for the test data using the final model
test_pred <- predict(final_model, newdata = pcr_test_scores)

# Plot Predicted vs Actual for Test Data
plot(pcr_test_scores$y, test_pred,
      main = "Predicted vs Actual Response (Test Data)",
      xlab = "Actual Response",
      ylab = "Predicted Values",
      pch = 16, col = "lightgreen")
abline(0, 1, col = "red", lwd = 2)

```



```

# Calculate RMSE for Test Data
test_rmse <- sqrt(mean((pcr_test_scores$y - test_pred)^2))
cat("Test RMSE:", test_rmse, "\n")

```

```
## Test RMSE: 0.3480036
```

The RMSE for the test data dropped from 0.76 in the original model to 0.34 after using the best subset model, by this we can conclude that the best subset model significantly improves prediction accuracy. This means the best subset model makes more accurate predictions because it focuses on the most important parts of the data.